

Introduction to Natural Language Processing

a.k.a NLP



What is going to happen today

An introduction


Together, building a spam classifier

Your turn, sentiment analysis

Together again, toward improving our spam classifier

A bit about me



Senior Data Scientist at  in Paris

(a **FREE**NOW  company)

Curious about NLP and happy to share
my recent learnings with you!

What if you fall in love with NLP and want to know more?

deeplearning.ai just launched a [new specialization, all about NLP](#), on Coursera

Browse > Data Science > Machine Learning

Offered By


Natural Language Processing Specialization




Break into the NLP space. Master cutting-edge NLP techniques through four hands-on courses!

★★★★★ 4.7 580 ratings

Why care about NLP?



Nice but what for?



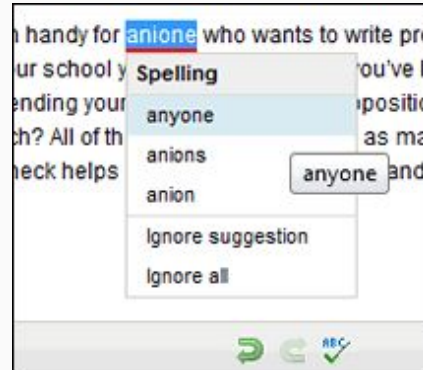
Machine translation



*Speech recognition +
Question answering*



Chatbots



Spell checking

How to work with text

How to work with text

Algorithms know well how to work with numbers

so how to with text = how to meaningfully transform text into numbers

Bag-of-words approach

Converting text into numbers

a.k.a Text preprocessing

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

"List listed lists listing listings."



['List', 'listed', 'lists', 'listing', 'listings', '.']

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

To “convert a collection of text documents to a matrix of token counts” [1]

[1] [Scikit-learn documentation on CountVectorizer](#)

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

An example (from [1])

```
corpus = [  
    'This is the first document.',  
    'This document is the second document.',  
    'And this is the third one.',  
    'Is this the first document?',  
]
```

```
tokens = ['and', 'document', 'first', 'is', 'one', 'second', 'the', 'third', 'this']
```

Vectorizer output:

```
[[0 1 1 1 0 0 1 0 1]  
 [0 2 0 1 0 1 1 0 1]  
 [1 0 0 1 1 0 1 1 1]  
 [0 1 1 1 0 0 1 0 1]]
```

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

Beware: it gives a lot of weights to frequent
(and maybe not so informative) words...

⇒ TF-IDF fixes this

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

TF-IDF: Term Frequency - Inverse Document Frequency

⇒ to measure how important a word is to a document in a corpus

ex: A frequent word in a document that is also frequent in the corpus is less important to a document than a frequent word in a document that is not frequent in the corpus

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

```
corpus = [  
    'This is the first document.',  
    'This document is the second document.',  
    'And this is the third one.',  
    'Is this the first document?',  
]
```

```
tokens = ['and', 'document', 'first', 'is', 'one', 'second', 'the', 'third', 'this']
```

Vectorizer output:

```
[[ 0, 0.5, 0.6, 0.4, 0, 0, 0.4, 0, 0.4]  
 [ 0, 0.7, 0, 0.3, 0, 0.5, 0.3, 0, 0.3]  
 [0.5, 0, 0, 0.3, 0.5, 0, 0.3, 0.5, 0.3]  
 [ 0, 0.5, 0.6, 0.4, 0, 0, 0.4, 0, 0.4]]
```

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

CountVectorizer and TF-IDF huge limitation

⇒ corpus dimension (*number of unique tokens*)

To reduce dimension:

Normalization

Stemming

Lemmatization

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

['List', 'listed', 'lists', 'listing', 'listings', '.']



['list', 'listed', 'lists', 'listing', 'listings', '.']

Underlying question: do we want to discriminate between “List” and “list”?

Sometimes we do:

⇒ “White House” versus “white house”

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

['list', 'listed', 'lists', 'listing', 'listings', '.']



['list', 'list', 'list', 'list', 'list', '.']

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

['list', 'listed', 'lists', 'listing', 'listings', '.']



['list', 'listed', 'list', 'listing', 'listing', '.']

Text preprocessing

Tokenization

CountVectorizer

TF-IDF

Normalization

Stemming

Lemmatization

Stemming or Lemmatization?

- It depends :)
- Stemming faster
- Lemmatization more informative

Text classification

You can use your favorite classifier with text

Naive Bayes or Logistic Regression usually provide a nice baseline

Performance assessed through AUC Score (Area Under the receiver operating characteristic Curve)

A word about the data we will use

SMS Spam Collection v. 1

“Public set of SMS labeled messages that have been collected for mobile phone spam research. [...] A collection of 5,574 English [...] messages.”

Dedicated web page: <http://www.dt.fee.unicamp.br/~tiago/smsspamcollection/>

Reference paper: Almeida, T.A., Gómez Hidalgo, J.M., Yamakami, A. **Contributions to the Study of SMS Spam Filtering: New Collection and Results**. Proceedings of the 2011 ACM Symposium on Document Engineering (DOCENG'11), Mountain View, CA, USA, 2011.

For sentiment analysis



 Dataset

Amazon Reviews: Unlocked Mobile Phones

More than 400,000 reviews from Amazon's unlocked mobile phone category

 PromptCloud • updated 4 years ago (Version 1)

Available on Kaggle platform

Let's move on to classifying spam

And now, sentiment analysis!